

UPLOAD SLOT 06 OF 12 · AI METHODOLOGY REPORT

AI Methodology Report

Four models, the discipline behind them, and what each one changed about the design

The challenge asks how AI supported the design. This submission's answer is a reproducible analysis pipeline that runs end to end with one command — not a description of prompting an image generator. Analysis that does not change a drawing is decoration; each model below is reported with the design decision it altered.

Al Safa 2 Park · **Falaj Al Safa** · 15,000 m² · Al Safa 2, Dubai · Mohamed Wasim, Individual Applicant · Issued 03 August 2026 · every quantity regenerated by python `run_analysis.py`

1. The discipline that makes these models mean anything


The target must not be recoverable from the inputs by algebra. It is easy to score $R^2 = 1.000$ on this project by accident: the heat index is a closed-form function of temperature and humidity, so 'predicting' it from temperature and humidity is arithmetic wearing a lab coat. Every model below was designed against that failure, and the pipeline asserts the absence of leaked features as a test.

2. The four models

Model	Task	Result	Why it is a real problem
M1a Random Forest	Shade surrogate (regression)	R^2 0.998	Learns a slow ray-traced simulation from cheap plan geometry
M1b Neural network	Shade surrogate (deployed)	R^2 0.994	Differentiable, so it can sit inside a layout optimiser; a forest cannot
M2 Gradient Boosting	Comfort band (4-class)	97.5%	Temperature and humidity withheld — it sees only sun position and the calendar
M3 K-Means	Microclimate regimes	$k=2$	Unsupervised; k is <i>selected</i> by silhouette score, not chosen to look tidy

3. Why M1 is the flagship


Ray-tracing annual shade means projecting 131 tree shadows onto 15,000 ground cells for every daylight hour. It is slow, and it must be re-run for every design variation. That cost is precisely what stops a designer exploring. The surrogate answers in milliseconds instead of minutes, at R^2 0.994 against ray-traced ground truth on a held-out test set. That is AI changing how the design was made, rather than decorating it afterwards.



Neural-network predictions against ray-traced ground truth on the held-out test set. Analysis output — computed from project data.

4. Why M2 answers a question worth asking

If thermal stress is predictable from a clock and an ephemeris alone, park operations can be scheduled without a sensor network — a real capital and maintenance saving over the life of the park. At 97.5% accuracy with temperature and humidity withheld, it is. Errors fall between adjacent comfort bands, never between comfortable and dangerous.



Confusion matrix, temperature and humidity withheld. Analysis output — computed from project data.

5. What the models actually changed

Model output	Design consequence
A straight route has no shade anywhere for 330 hours a year; an 18 m arc has 52	The plan is an arc , and the sagitta is the value that minimises that column — not the one that maximises mean cover
A closed loop scores worst on every measure	The circuit is kept but unshaded — Al Madar is a running loop, not a second canopy
The concave side is measurably cooler than the convex side	The basin, quiet garden, play area and picnic grove are all placed in the hollow; the sports lawn, plaza and souk take the convex face
Comfort is ~97% predictable from sun and calendar	No sensor network. Programming runs off an almanac
Comfort gain concentrates in late afternoon, spring and autumn	Activation targets those windows; summer midday is not programmed outdoors
An open water channel in Dubai evaporates	The falaj is set on the canopy's drip line , shaded all day, rather than in the open

6. Where AI was deliberately not used

Visitor demand is a scenario model, not a machine learning result. No footfall data exists for this site. A model trained on invented demand would only recover the assumption that produced it, and reporting that as a finding would be dishonest. It is therefore excluded from the model suite and labelled a scenario wherever it appears.

The photoreal renders are illustrations, not analysis. They are captioned as artistic impressions throughout. Three earlier visuals that presented invented data as measurement were withdrawn entirely.

7. Reproducibility

The complete analysis — data, code, trained models and tests — runs end to end with `python run_analysis.py` and is verified by `python -m tests.test_pipeline` (38 checks), `node docs/_PORTAL/selftest.js` (64 checks) and a frame-by-frame test of the concept film. A juror who wants to know where a number comes from can be given a file path rather than an opinion.

The full repository — every dataset, every model, every line of the pipeline above — is published at https://github.com/wasim437/Al_SAFA and the live analytics portal at https://wasim437.github.io/Al_SAFA/. Both run from the same code that produced this document.